

TITLE

DAN SZYMCZAK, McMaster University, Canada
JACQUES CARETTE and SPENCER SMITH

CONTEXT: Software (re-)certification requires the creation and maintenance of many different software artifacts. Manually creating and maintaining them is tedious and costly.

OBJECTIVE: Improve software (re-)certification efforts by automating as much of the artifact creation process as possible while maintaining full traceability within – and between – artifacts.

METHOD: Start by analyzing the artifacts themselves from several case studies to understand what (semantically) is being said in each. Capture the underlying knowledge and apply transformations to create each of the requisite artifacts through a generative approach.

RESULTS: Case studies – GlassBR to show capture and transformation. SWHS and NoPCM for reuse (Something about Kolmogorov complexity / MDL here?). Captured knowledge can be re-used across projects as it represents the “science”. Maintenance involves updating the captured knowledge or transformations as necessary. Creation of our tool – Drasil – facilitates this automation process using a knowledge-based approach to Software Engineering.

CONCLUSIONS: With good tool support and a front-loaded time investment, we can automate the generation of software artifacts required for certification. (fill in later)????

Additional Key Words and Phrases: ??

ACM Reference Format:

Dan Szymczak, Jacques Carette, and Spencer Smith. 2018. TITLE. *ACM Trans. Softw. Eng. Methodol.* 1, 1 (April 2018), 8 pages. <https://doi.org/0000001.0000001>

1 INTRODUCTION

Writing non-executable software artifacts (requirements and design documents, verification & validation plans, etc.) can be tedious work, but is ultimately necessary when attempting to certify software. Similarly, maintenance of these artifacts, as necessary for re-certification as improvements are made, typically requires a large time investment.

Why, in a world of software tools, do we continue to undertake these efforts manually? Literate programming had the right idea, but was too heavily focused on code.

We want to aid software (re-)certification efforts by automating as much of the artifact creation process as possible. By generating our software artifacts – including code – in the right way, we can implement changes much more quickly and easily for a modest up-front time investment. By front-loading the costs of maintenance and rolling them into the development cycle, we can save time and money in the long run.

Authors' addresses: Dan Szymczak, McMaster University, 1280 Main St. W., Hamilton, ON, L8S 4K1, Canada, szymczdm@mcmaster.ca; Jacques Carette; Spencer Smith.

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for components of this work owned by others than the author(s) must be honored. Abstracting with credit is permitted. To copy otherwise, to republish, to post on servers or to redistribute to lists, requires prior specific permission and/or a fee. Request permissions from permissions@acm.org.

© 2018 Copyright held by the owner/author(s). Publication rights licensed to the Association for Computing Machinery. 1049-331X/2018/4-ART \$15.00

<https://doi.org/0000001.0000001>

1.1 Software (Re-)certification

When we talk about software certification, we are specifically discussing the goal of determining “based on the principles of science, engineering and measurement theory, whether an artifact satisfies accepted, well defined and measurable criteria” [6]. Essentially, we are ensuring that software, or some piece of it, performs a given task to within an acceptable standard and can potentially be reused in other systems.

Software certification is necessary by law in certain fields. This is particularly evident in safety-critical applications, such as control systems for nuclear power plants or X-ray machines.

Different certifying bodies exist across domains and each has their own list of requirements to satisfy for certifying software. Looking at some examples [1–3, 8] there are many pieces of requisite documentation including, but not limited to:

- Problem definition
- Theory manual
- Requirements specification
- Design description
- Verification and Validation (V&V) report

We should keep in mind that we require full traceability – inter- and intra-artifact – of the knowledge contained within these artifacts. That is, we should be able to find an explicit link between our problem definition and theory manual, down to our requirements, design, and other development planning artifacts. From there, we should be able to continue through our proposed verification and validation plans, and should eventually end up in the V&V report.

Ensuring this traceability and, in fact, getting anything certified has many costs associated with it. There is a massive time investment, fees, and costs associated with contracting out a third-party verifier. Overall it is a very expensive process.

Re-certification of software following any change, no matter how minor, incurs a similar level of costs; all the artifacts must be updated to reflect the new change, and everything must be re-checked and verified to ensure no new errors have been introduced.

We intend to alleviate some of this cost-burden through a strategic, generative approach to Software Engineering (SE). With the automated generation of artifacts, we can implement changes quickly and automatically update every relevant and/or dependent artifact.

1.2 Document-Driven Design

- Document-driven design and its advantages/disadvantages. - Quality improvements - Large time investment for initial artifact creation - Large time investment for every (non-trivial) update - Often artifacts fall out of sync with each other and are inconsistent - All but necessary for certification - Need to overcome the disadvantages somehow. - Most obvious solution would be to automate if possible.

1.3 Scope

- Scientific Computing Software - Why? Many highly specialized SCS require certification. Ex. Control sfwr in nuclear power, x-ray machines, and other safety-critical contexts. - Well understood domain -> theories underpinning the work being done.

2 BACKGROUND

- We are not the first to try and deal with certification / artifact creation.

Fig. 1. Data Definition for !FIXME! from GlassBR SRS

Fig. 2. Data Definition code from GlassBR implementation

2.1 Previous efforts

Previous attempts at automating / reducing the artifact burden.

- Compendia - Trying to solve the problem of reproducibility - Fits with goals of certification - focused on good science and being able to re-run experiments exactly - Not focused on DDD or its benefits, moreso

- Previous attempts at automatically generating documentation - LP, tools like javadoc, Haddock, etc. - Too code-centric! - Comments and code still need to be updated in parallel, albeit to a lesser extent in some cases - In general, fairly rigidly structured output (you don't have much say on how it looks, only what information should be included and, sometimes, where - Finish with a focus on the good stuff: - Identified the need for good documentation - Keeps docs and code in the same place - Easier to manually maintain consistency and apply updates - One other problem we've identified: - common underlying knowledge between projects is duplicated as there is no real cross-project reuse mechanism in place with these tools.

3 A RATIONAL ANALYSIS OF SOFTWARE ARTIFACTS

- This section exists to show how we get from problem to solution. - We introduce our case studies in a bit more depth here - GlassBR - what it's for, if it'll - SWHS and NoPCM - Program family members with a twist. - The rest (tiny, Gamephysics, and SSP) for additional examples and to give us a bit more credibility in our analysis. - Looking for commonalities between types of artifacts and what they are really saying. - An obvious commonality across many projects in SCS - SI and derived Units.

3.1 Common software artifacts

- Compare and contrast different software artifacts. - SRS vs. detailed design vs. code - same knowledge, different 'views' - only some of that knowledge is necessarily relevant in those views - **!FIGURE!:** SRS & DD showing the same piece of knowledge in diff contexts. Use a few different **!FIGURE!:**here. - **!FIGURE!:** Attempt to show generalized overlap via Venn diagram?

3.2 Emerging structures

- As shown above,

In the common software artifacts we see different ways of representing what are, semantically, the same things (for example, see Figure 3.1). We are really seeing the pieces of underlying knowledge that have been composed from a variety of components. Each component tells us something about one aspect of that piece of knowledge. Particularly, they give examples of how we can transform, or view, the same semantic knowledge in different contexts.

If we take a look at one particular example across artifacts from GlassBR (Figures 1,2), we can see that it is an aggregation of the following components:

- Unique Identifier (label)
- Symbolic (theory) representation
- Symbolic (implementation) representation
- Concise natural language description (a term)

Table 1. Knowledge Classes

Knowledge Class	ID	Term	Abbreviation	Definition	Symbol	Equation	Constraints	Units
Labeled	X							
Named Idea	X	X	O					
Common Idea	X	X	X					
Concept	X	X	X	X				

Legend: X - Mandatory; O - Optional

- Verbose natural language description (a definition)
- Equation
- Constraints
- Units?

The unique identifier is fairly straightforward (!FIXME id!), it is just a label that we associate with this particular piece of knowledge and nothing else. The symbolic representations are just the symbols we use when referring to this particular quantity in an equation (theory) or code (implementation) context. Our natural language descriptions are terms and their corresponding definitions (!FIXME! and !FIXME! respectively for this example).

We also have a defining equation, which incorporates the symbolic representation for various other pieces of knowledge and relates them to !FIXME name!. Similarly, we have constraints which are just relationships which must be maintained between !NAME! and some other quantities. Lastly, we have the units which our quantity is measured in, which are derived from the fundamental !SI UNITS!.

Similar examples of knowledge crop up over all the artifacts. Some have the same depth of information, whereas others do not. Regardless, all of our knowledge shares some components in common. We will always have a label, and usually a term and definition. Depending on what we’re looking at, there may not be a symbolic representation, or perhaps we have a quantity that is unit-less. These special cases help us see the underlying root structure from which our knowledge buds.

-Discuss the breakdown of knowledge into classes. Refer to Table 1 for more.

4 KNOWLEDGE-BASED SOFTWARE ENGINEERING (KBSE)

Knowledge-Based Software Engineering (KBSE) was originally defined as an “engineering discipline that includes the integration of knowledge into software systems in order to solve complex problems, which would normally require rather high level of human expertise” [5]. This is a solid definition, provided we understand what “knowledge” is. So then, what exactly is knowledge?

Knowledge “presents understanding of a subject area. It includes concepts and facts ... as well as relations ... and mechanisms for how to combine them to solve problems in that area” [4].

For our purposes, we extend and tighten this definition to include the additional constraint that a piece of knowledge has a structured encoding, as opposed to natural language encoding, which then allows it to be automatically reused. For example, the first law of thermodynamics is a piece of knowledge that can be simply expressed as “total energy within a closed system must be conserved”,

but this is not a structured encoding. One such encoding would allow us to view the knowledge in those simple terms, or just as easily, we could view it as:

$$\Delta U = Q - W$$

Regardless of our view, the underlying structured knowledge encoding does not change – we are merely projecting out what is relevant to our current audience.

For our KBSE approach to succeed, there are two major requirements. First off, we must capture the underlying knowledge in a meaningful way that can be reused across artifacts. We want a single source for our knowledge, regardless where it ends up or how it is viewed. This allows us, using the right transformations, to automatically generate our software artifacts from the underlying knowledge-base.

The second requirement is that we restrict our scope to well-understood domains as we need a solid theoretical underpinning. Both mathematics and the physical sciences are good examples of well-understood domains as the knowledge has already been formalized and, to an extent, structured. These are also good candidate domains since we need to explain the underlying knowledge to a computer in a nontrivial way, which from our experience, is harder than it sounds.

With that in mind, we have decided to restrict our focus to KBSE for Scientific Computing Software (SCS) as it is a field rich in knowledge we can use.

4.1 Capturing Knowledge

From our work in Section 3.2 we can create a knowledge-capture mechanism for encoding the requisite underlying science into a machine-usable form. By laying out the structure, we can see which information must be captured for each piece of knowledge.

Different types of information are required for encoding each of the various pieces of knowledge we intend to use. Some types of knowledge lack specific information bindings, for example a *named idea* does not necessarily have a symbol associated with it, however, a *quantity must* have a symbol alongside its *term* – the fundamental information in a named idea.

We borrow the idea of *Chunks* from Literate Programming (LP) [7] to facilitate our knowledge-capture. A chunk in its most rudimentary sense is simply a labeled piece of information. Given our understanding of how the knowledge should be structured, we have created a hierarchy of classes built up from the simplest of chunks, to fulfill our knowledge-capture requirements. This hierarchy as implemented in Drasil can be seen in Figure 3. It mimics the structure mentioned in Section 3.2. We will delve deeper into the specifics of our hierarchy in Section 5.

- We have to capture all of the information surrounding a piece of knowledge to create our artifacts, regardless of whether that information is relevant to any one particular artifact. - Once knowledge is properly captured, we shouldn't have to capture it again if we want to reuse it in a different project.

4.2 (Re-)Using Knowledge

- Most obvious benefit -> no more copy/paste! Just reuse the chunk you need. - Transformations - Represent different 'views' of the knowledge based on how abstract, what audience, etc. - Translate the knowledge into its requisite form (eqns, descriptions, code) - Variabilities -> different projects in the same family. - Easy to specialize to different family members - Example: SWHS vs NoPCM **!FIGURE!** Show portion of each SRS, one similarity, one difference? - Requires a framework / tool support to automate rendering of these transformations, otherwise it is even more work for humans.

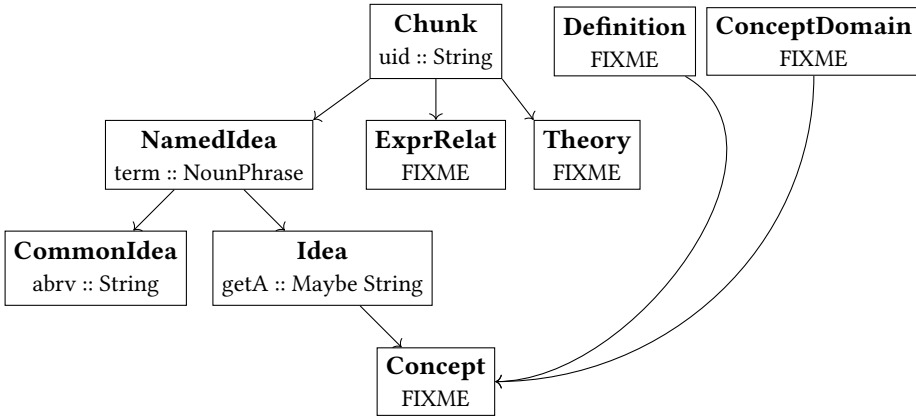


Fig. 3. Chunk hierarchy in Drasil Today

5 DRASIL

- To use KBSE to its potential we need a strong support framework - Intro to Drasil **!FIGURE!** Knowledge tree - What it is and does - Domain Specific Language - Generate all the things! - Dev to date. - How is Knowledge Capture handled in Drasil? - chunks! - What do transformations look like? Recipes! **!FIGURE!** SmithEtAl template for SRS = Drasil.DocumentLanguage - Key components of the generator / renderer

5.1 Developing Drasil - A grounded theory

- Following grounded theory (ish). Using data from case studies to guide development and implement new features. - **!FIGURE!**: Before and after System Information. - **!FIGURE!**: Before and after mini-DBs - Majority of features developed after analyzing commonalities in the case studies and abstracting them out. - Allows for rapid progress -> constant iteration based on what we find in the data.

5.2 Drasil Today

- Sentence and Document - Explain the chunk hierarchy (refer to Section 4.1 figure) - Data.Drasil **!FIGURE!** Knowledge areas we've started to capture (See: SE-CSE paper) - Recipe Language(s) - Refer to: **!FIGURE!** Drasil.DocumentLanguage - The generator - HTML and TeX rendering - GOOL for code - System Information -> Get into it

6 CASE STUDIES - IN MORE DEPTH

- Re-introduce case studies - Our methods for reimplementing - CI for testing - Start showing off re-use and automated generation. - Start with common knowledge (generalized **!FIGURE!?**) - Then onto GlassBR example to show off the doc lang recipe (**!FIGURE!?**) - Then let's see SRS vs. NoPCM for reuse (particularly NoPCM) (**!FIGURE!?**)

6.1 Data.Drasil

- Common knowledge **!FIGURE!** SI_Units **!FIGURE!** Thermodynamics (ConsThermE?)

6.2 GlassBR

- Brief intro to problem GlassBR is solving - how it works - Show off the doc language here **!FIGURE!** GlassBR SRS in (truncated) DocLang format - "Reads like a table of contents, with a few quirks" - Show off some code generation **!FIGURE!** Side-by-side of Chunk Eqn vs. Doc Eqn vs. Code - "Easy to see that the code matches the equations" - Talk about potential variabilities and how to make this a family - Why is this interesting? - Fairly straightforward example of something a scientist would create/use in their research

6.3 NoPCM & SWHS

- Re-introduce the problems - See how they're a family? - Really drill in the similarities **!FIGURE!** Figure showing NoPCM import(s) - Lots of knowledge-reuse - Very few 'new' chunks (count them?) - Show example of variability in action **!FIGURE!** Equation with/without PCM (rendered?) - Why this example is interesting: - ODE solver -> We don't gen, just link to existing good one(s)

6.4 Others

- Mention SSP, Tiny, GamePhysics, but don't go too in-depth. - Useful examples as they give us a wider range of problems for analysis - Testing - Physics is physics -> when we make updates, the underlying knowledge isn't changing, so neither should our output - Refer to CI

6.5 Freebies - Compliments of System Information

- Thanks to the recipe language and the way we structure out system information we can get - Table of Symbols - Table of Units - Table of Abbreviations and Acronyms - Bibliography
- All tedious to do by hand, but are free to automatically generate - Generator includes sanity-checking -> Can't use something that isn't defined! - Sanity-checks are 'free' -> we can check for errors with our symbols, ensure units are consistent, guard against constraints, and ensure we only reference those things which are defined in our system. - Sanity-checks are run every time artifacts are generated.

6.6 Results

- Here we discuss the results we've seen so far. - Had some of these case studies attempted to be certified, they would (should) have failed. - A number of common problems.

6.7 Common issues across case studies

- A number of undefined symbols even after multiple passes by humans. (Auto-generating the symbol table and including sanity-checking revealed them)

6.8 NoPCM and SWHS

- Along with the common errors, there was some sharing of PCM-related knowledge - Found because PCM symbols were not in the ToS and the sanity-check caught it. - No way to specifically exclude knowledge that shouldn't 'exist' in a project - Work in Kolmogorov complexity / MDL for NoPCM + SWHS? - Kolmogorov/MDL implies less writing for the same artifacts -> less to sift through = maybe better?

6.9 SSP

- Symbols for given quantities changed throughout the documentation - Went unnoticed by a human for years! Found almost instantly by Drasil - the new symbols were undefined.

6.10 Pervasive Bugs

- Mistakes in knowledge can be found in all artifacts - more likely to be caught! - Easy to track down errors (smart error messages point to the exact chunk causing the problem).

7 FUTURE WORK

[*SS* - Once we are capable of true variability in the documentation, we can really start asking the question about what is the "best" documentation for a given context. In the future experiments could be done with presenting the same information in different ways to find which approach is the most effective.]

[*SS* - Related to the previous point, the act of formalizing the knowledge that goes into the requirements documentation forces us to deeply understand the distinctions between difference concepts, like scope, goal, theory, assumption, simplification, etc. With this knowledge we can improve the focus and effectiveness of existing templates, and existing requirements solicitation and analysis efforts. Teaching it to a computer.]

- Run an experiment to determine how easy it is to create new software with Drasil.
- Run an experiment to see how easy it is to find and remove errors with Drasil
- Experiment to see time saved in maintenance while using Drasil vs. not

8 CONCLUSION

- Easier to find errors (anecdotally) - future work will tell us if this holds.

REFERENCES

- [1] CENTER FOR DEVICES AND RADIOLOGICAL HEALTH, CDRH. General principles of software validation; final guidance for industry and FDA staff. Tech. rep., US Department Of Health and Human Services Food and Drug Administration Center for Devices and Radiological Health Center for Biologics Evaluation and Research, York, England, January 2002.
- [2] CSA. Quality assurance of analytical, scientific, and design computer programs for nuclear power plants. Tech. Rep. N286.7-99, Canadian Standards Association, 178 Rexdale Blvd. Etobicoke, Ontario, Canada M9W 1R3, 1999.
- [3] CSA. Guideline for the application of N286.7-99, quality assurance of analytical, scientific, and design computer programs for nuclear power plants. Tech. Rep. N286.7.1-09, Canadian Standards Association, 5060 Spectrum Way, Suite 100, Mississauga, Ontario, Canada L4W 5N6, 1-800-463-6727, 2009.
- [4] DURKIN, J. *Expert systems: design and development*. Macmillan Coll Div, 1994.
- [5] FEIGENBAUM, E. A., AND MCCORDUCK, P. *The fifth generation*. Addison-Wesley, 1983.
- [6] HATCLIFF, J., HEIMDAHL, M., LAWFOR, M., MAIBAUM, T., WASSYNG, A., AND WURDEN, F. A software certification consortium and its top 9 hurdles. *Electronic Notes in Theoretical Computer Science* 238, 4 (2009), 11–17.
- [7] KNUTH, D. E. Literate programming. *The Computer Journal* 27, 2 (1984), 97–111.
- [8] U.S. FOOD AND DRUG ADMINISTRATION. Infusion pumps total product life cycle: Guidance for industry and fda staff. on-line, December 2014.