

GOOL: A Generic Object-Oriented Language

Anonymous Author(s)

Abstract

Text of abstract

Keywords keyword1, keyword2, keyword3

1 Introduction

Given a task, before writing any code a programmer must select a programming language to use. Whatever they may base their choice upon, almost any programming language will work. While a program may be more difficult to express in one language over another, it should at least be possible to write the program in either language. Just as the same sentence can be translated to any spoken language, the same program can be written in any programming language. Though they will accomplish the same tasks, the expressions of a program in different programming languages can appear substantially different due to the unique syntax of each language. Within a single language paradigm, such as object-oriented (OO), these differences should not be as extreme – at least the global structuring mechanisms and the local idioms will be shared. Mainstream OO languages generally contain (mutable) variables, methods, classes, objects and a core imperative set of primitives. Some OO languages even have very similar syntax (such as Java and C# say).

When faced with the task to write a program meant to fit into multiple existing infrastructure, which might be written in different languages, frequently that entails writing different versions of the program, one for each. While not necessarily difficult, it nevertheless requires investing the time to learn the idiosyncrasies of each language and pay attention to the operational details where languages differ. Ultimately, the code will likely be marred by influences of the language the programmer knows best. They may consistently use techniques that they are familiar with from one language, while unaware that the language in which they are currently writing offers a better or cleaner way of doing the same task [5, 18]. Besides this likelihood of writing sub-optimal code, repeatedly writing the same program in different languages is entirely inefficient, both as an up-front development cost, and even more so for maintenance.

Since languages from the same paradigm share many semantic similarities, it is tempting to try to leverage this; perhaps the program could be written in one language and automatically translated to the others? But a direct translation is often difficult, as different languages require the programmer to provide different levels of information, even to achieve the same tasks. For example, a dynamically typed

language like Python cannot be straightforwardly translated to a statically typed language like Java, as additional type information generally needs to be provided¹.

What if, instead, there was a single meta-language which was designed to contain the common semantic concepts of a number of OO languages, encoded in such a way that all the necessary information for translation was always present? This source language could be made to be agnostic about what eventual target language was used – free of the idiosyncratic details of any given language. This would be quite the boon for the translator. In fact, we could try to go even further, and attempt to teach the translator about idiomatic patterns of each target language.

Why would this even be possible? There are commonly performed tasks and patterns of OO solutions, from idioms to architecture patterns, as outlined in [10]. A meta-language that provided abstractions for these tasks and patterns would make the process of writing OO code even easier.

But is this even feasible? In some sense, this is already old hat: most modern compilers have a single internal Intermediate Representation (IR) which is used to target multiple processors. Compilers can generate human-readable symbolic assembly code for a large family of CPUs. But this is not quite the same as generating human-readable, idiomatic high-level languages.

There is another area where something like this has been looked at: the production of high-level code from Domain-Specific Languages (DSL). A DSL is a high-level programming language with syntax and semantics tailored to a specific domain [16]. DSLs allow domain experts to write code without having to concern themselves with the details of General-Purpose programming Languages (GPL). A DSL abstracts over the details of the code, providing notation for a user to specify domain-specific knowledge in a natural manner. Such DSL code is typically translated to a GPL for execution. Abstracting over code details and compiling into traditional OO languages is exactly what we want to do! The details to abstract over include both syntactic and operational details of any specific language, but also higher-level idioms in common use. Thus the language we are looking for is just a DSL in the domain of OO programming languages!

There are some DSLs that already generate code in multiple languages, to be further discussed in Section 6, but none of them have the combination of features we want. We are indeed trying to do something odd: writing a “DSL” for what is essentially the domain of OO GPLs. Furthermore, we have additional requirements:

PL’18, January 01–03, 2018, New York, NY, USA
2018.

¹Type inference for Python notwithstanding

1. The generated code should be human-readable,
2. The generated code should be idiomatic,
3. The generated code should be documented,
4. The generator should allow one to express common OO patterns.

We have developed a Generic Object-Oriented Language (GOOL), demonstrating that all these requirements can be met. GOOL is a DSL embedded in Haskell that can currently generate code in Python, Java, C#, and C++². Others could be added, with the implementation effort being commensurate to their (semantic) distance to the languages already supported.

First we present the high-level requirements for such an endeavour, in Section 2. To be able to give illustrated examples, we next show the syntax of GOOL in Section ?? The details of the implementations, namely the internal representation and the family of pretty-printers, is in Section 4. Common patterns are illustrated in Section 5. We close with a discussion of related work in Section 6, plans for future improvements in Section ??, and conclusions in Section 8.

2 Requirements

While we outlined some of our requirements above, here we will give a complete list, along with acronyms (to make referring to them simpler), as well as some reasoning behind each requirement.

- mainstream** Generate code in mainstream object-oriented languages.
- readable** The generated code should be human-readable,
- idiomatic** The generated code should be idiomatic,
- documented** The generated code should be documented,
- patterns** The generator should allow one to express common OO patterns.
- common** Language commonalities should be abstracted.

Targetting OO languages (**mainstream**) is primarily because of their popularity, and thus would enjoy the most potential users – in much the same way that the makers of Scala and Kotlin chose to target the JVM to leverage the Java ecosystem, and Typescript for Javascript.

The **readable** requirement is not as obvious. As DSL users are typically domain experts who are not “programmers”, why generate readable code? Few Java programmers ever look at JVM bytecode, and few C++ programmers look at assembly. But GOOL’s aim is different: to allow writing high-level OO code once, but have it be available in many GPLs. One use case would be to generate libraries of utilities for a narrow domain. As needs evolve and language popularity changes, it is useful to have it immediately available in a number of languages. Another use, which is a core part of our own motivation, is to have *extremely well documented* code, indeed to a level that would be unrealistic to do by

²and is close to generating Lua and Objective-C, but those backends have fallen into disuse

hand. But this documentation is crucial in domains where *certification* of code is required.

The same underlying reasons for **readable** also drive **idiomatic** and **documented**, as they contribute to the human-understandability of the generated code. **idiomatic** is important as many human readers would find the code “foreign” otherwise, and would not be keen on using it. Note that documentation can span from informal comments meant for humans, to formal, structured comments useful for generating API documentation with tools like Doxygen, or with a variety of static analysis tools. Readability (and thus understandability) are improved when code is pretty-printed[7]. Thus taking care of layout, redundant parentheses, well-chosen variable names, using a common style with lines which are not too long, are just as valid for generated code as for human-written code. GOOL does not prevent users from writing undocumented or complex code, if they choose to do so. It just makes it easy to have **readable**, **idiomatic** and **documented** code in multiple languages.

The **patterns** requirement is typical of DSLs: common programming patterns can be reified into a proper linguistic form instead of being merely informal. In particular some of the *design patterns* of [10] can become part of the language itself. This does make writing some OO code even easier in GOOL than in GPLs, it also helps quite a lot with keeping GOOL language-agnostic and generating idiomatic code. Illustrative examples will be given in Section 5. But we can give an indication now as to why this helps: Consider Python’s ability to return multiple values with a single return statement, which is uncommon in other languages. Two choices might be to disallow this feature in GOOL, or throw an error on use when generating code in languages that do not support this feature. In the first case, this would likely mean unidiomatic Python code, or increased complexity in the Python generator to infer that pattern. The second option is worse still: one might have to resort to writing language-specific GOOL, obviating the whole reason for the language! Multiple-value return statements are always used when a function returns multiple outputs; what we can do in GOOL is to support such multiple-value return, and then generate the idiomatic pattern of implementation in each target language.

Our last requirement, that language commonalities (**common**) be abstracted is an internal requirement: we noticed a lot of repeated code in our backends, something that ought to be distasteful to most programmers. For example, writing a generator for both Java and C^{sharp} makes it incredibly clear how similar the two languages are.

3 Creating GOOL

How do we go about creating a “generic” object-oriented language? We chose an incremental abstraction approach: start from two languages, and unify them *conceptually*. In

other words, pay very close attention to the *denotational* semantics of the features, some attention to the operational semantics, and ignore syntactic details.

This is most easily done from the core imperative language outwards. Most languages provide similar basic types (variations on integers, floating point numbers, characters, strings, etc) and functions to deal with them. The core expression language tends to be extremely similar cross languages. One then moves up to the statement language — assignments, conditionals, loops, etc. Here we start to encounter variations, and choices can be made, and we'll cover that later.

For ease of experimentation, we chose to make GOOL an embedded domain specific language (EDSL) inside Haskell. Haskell is very well-suited for this task, offering a variety of features (GADTs, type classes, parametric polymorphism, kind polymorphism, etc) which is extremely useful for building languages. Its syntax is also fairly liberal, so that it is possible to create *smart constructors* that somewhat mimic the usual syntax of OO languages.

3.1 GOOL Syntax: Imperative core

As our exposition has been somewhat abstract until now, it is useful to dive in and give some concrete syntax, so as to be able to illustrate our ideas with valid code.

Specifically, Basic types in GOOL are `bool` for Booleans, `int` for integers, `float` for doubles, `char` for characters, `string` for strings, `infile` for a file in read mode, and `outfile` for a file in write mode. Lists can be specified with `listType`. For example, `listType int` specifies a list of integers. Types of objects are specified using `obj` followed by the class name, so `obj "FooClass"` is the type of an object of a class called "FooClass".

Variables are specified with `var` followed by the variable name and type. For example, `var "ages" (listType int)` represents a variable called "ages" that is a list of integers. This illustrates a (necessary) design decision: even though we target languages like Python, as we also target Java, types are necessary. As type inference for OO languages is too difficult, we chose to be explicitly typed.

As some constructions are common, it is useful to offer shortcuts for defining them; for example, the above can also be done via `listVar "ages" int`. Typical use would be

```
let ages = listVar "ages" int in
```

so that `ages` can be used directly from then on.

Other keywords for specifying variables include `extVar` for variables from external libraries, `classVar` for accessing variables from a class, `objVar` for accessing variables from an object, and `self` for referring to an object in its own class definition, equivalent to `self` in Python or `this` in Java. The infix operator `$->` is an alternative to the `objVar` keyword.

Note that GOOL distinguishes a variable from its value³. To get the value of `ages`, one must write `valueOf ages`. The reason for this distinction will be made clear in section ??, driven by semantic considerations. This is beneficial for stricter typing and enables convenient syntax for **patterns** that translate to more idiomatic code.

all the following should be in two tables.

Literal values can be referred to by `litTrue`, `litFalse`, `litInt`, `litFloat`, `litChar`, and `litString`. Similar to those seen in most programming languages, GOOL provides many unary prefix operators and binary infix operators for defining expressions. In GOOL, each operator is prefixed with an additional symbol based on type. Operators that return Booleans are prefixed by a `?`, for example `?!` is used for negation, `?&&` for conjunction, `?||` for disjunction, and `?==` for equality. Operators on numeric values are prefixed by `#`, such as `#~` for negation, `#/^` for square root, `#|` for absolute value, `##` for modulus, and `#^` for exponentiation. Any other operators are prefixed by `$`, such as the previously mentioned `$->` operator for accessing a variable from an object.

another table. I'll put the contents in a list

- `inlineIf` conditional expression
- `funcApp` function application, to a list of parameters
- `extFuncApp` function application, for external library functions
- `newObj` for calling an Object constructor (`extNewObj` exists too)
- `objMethodCall` for calling a method on an object.

`selfFuncApp` and `objMethodCallNoParams` are two shortcuts for the common cases when a method is being called on `self` or when the method takes no parameters.

Variable declarations are statements, and take a variable specification as argument. For `foo = var "foo" int`, the corresponding variable declaration would be `varDec foo`, and to also initialize it `varDecDef foo (litInt 5)` can be used.

Assignments are represented by `assign a (litInt 5)`. Convenient infix and postfix operators are also provided, prefixed by `&`: `&=` is a synonym for `assign`, and C-like `&+=`, `&+=`, `&-=` and `&~` (the more intuitive `&--` cannot be used as `--` starts a comment in Haskell).

Other simple statements in GOOL include `break` and `continue`, `returnState` followed by a value to return, `throw` followed by an error message to throw, `free` followed by a variable to free from memory, and `comment` followed by a string to be displayed as a single-line comment.

Most languages have statement blocks, introduced by block with a list of statements in GOOL. Bodies (body) are composed of a list of blocks, and can be used as a function body, conditional body, loop body, etc. The purpose of blocks as an intermediate between statement and body is to allow

³as befits the use-mention distinction from analytic philosophy

for more organized, readable generated code. For example, the generator can choose to insert a blank line between blocks so lines of code related to the same task are visually grouped together. Naturally shortcuts are provided for single-block bodies (`bodyStatements`) and for the common single-statement case, `oneLiner`.

GOOL has two forms of conditionals: `if-then-else` via `ifCond` (which takes a list of pairs of conditions and bodies) and `if-then` via `ifNoElse`. For example:

```
ifCond [
  (foo ?> litInt 0, oneLiner (
    printStrLn "foo is positive")),
  (foo ?< litInt 0, oneLiner (
    printStrLn "foo is negative"))]
(oneLiner \$ printStrLn "foo is zero")
```

GOOL also supports `switch` statements.

There are a variety of loops. A `for-loop` (`for`) is parametrized by a statement to initialize the loop variable, a condition, a statement to update the loop variable, and a body; a `forRange` loop, which are given a starting value, ending value, and step size, as well as `forEach` loops. For example:

```
for (varDecDef age (litInt 0))
  (age < litInt 10) (age &++) loopBody
forRange age (litInt 0) (litInt 9)
  (litInt 1) loopBody
forEach age ages loopBody
```

While-loops (`while`) are parametrized by a condition and a body. Finally, `try-catch` statements (`tryCatch`) are parametrized by two bodies.

3.2 GOOL Syntax: OO features

A function declaration is followed by the function name, scope, binding type (static or dynamic), type, list of parameters, and body. Methods (`method`) are defined similarly, with the addition of the specification of the containing class' name. Parameters are built from variables, using `param` or `pointerParam`. For example, assuming variables “`num1`” and “`num2`” have been defined, one can define an `add` function as follows:

```
function "add" public dynamic_ int
  [param num1, param num2]
  (oneLiner (returnState (num1 #+ num2)))
```

The `pubMethod` and `privMethod` shortcuts are useful for public dynamic and private dynamic methods, respectively. `mainFunction` followed by a body defines the main function of a program. `docFunc` generates a documented function from a function description and a list of parameter descriptions, an optional description of the return value, and the function itself. This generate Doxygen-style comments.

Classes are defined with `buildClass` followed by the class name, name of the parent class (if applicable), scope, list of state variables, and list of methods. State variables can be built by `stateVar` followed by an integer, scope, static or dynamic binding, and the variable itself. The integer is a measure of delete priority. `constVar` can be used for constant state variables. Shortcuts for state variables include `privMVar` for private dynamic, `pubMVar` for public dynamic, and `pubGVar` for public static variables. For example:

```
buildClass "FooClass" Nothing public
  [pubMVar 0 var1, privMVar 0 var2]
  [mth1, mth2]
```

Nothing here indicates that this class does not have a parent, `privClass` and `pubClass` are shortcuts for private and public classes, respectively. `docClass` serves a similar purpose as `docFunc`.

3.3 GOOL syntax: modules and programs

Akin to Java packages and other similar constructs, GOOL has modules (`buildModule`) consisting of a module name, a list of libraries to import, a list of functions, and a list of classes. Module-level comments are done with `docMod`.

Finally, at the top of the GOOL hierarchy are programs, auxiliary files, and packages. A program is constructed by passing a program name and list of files to `prog`. Then a program and a list of auxiliary files can be passed to `package` to create a complete package, specified in GOOL. Auxiliary files are non-OO code files that augment the OO program. For example, a Doxygen configuration file is an auxiliary file that can be specified using the `doxConfig` keyword, and a makefile is another auxiliary file that can be specified with the `makefile` keyword. One of the parameters to the `makefile` keyword will toggle generation of a `make doc` rule, which will compile the Doxygen documentation with the generated Doxygen configuration file.

4 GOOL Implementation

GOOL is embedded in Haskell in the finally-tagless style originally described in [8]. In finally-tagless style, the internal representation of an embedded DSL is a group of functions defined in the host language, rather than the more traditional use of data constructors to build an interpretable abstract syntax tree (AST) for the DSL. Whereas an AST of data constructors must have a separately-defined interpreter, the functions themselves in finally-tagless compose to form the interpreter; the definitions of the functions describe how each is interpreted. Thus, the GOOL “keywords” referred to in the previous section were really Haskell functions that, when resolved, yield representations of the code to be generated. The finally-tagless style grants us more flexibility in the code we can generate and better extensibility since we do not need to pattern-match on AST nodes.

Finally-tagless facilitates development of a family of interpreters for a DSL by having the functions act on representations, where a type class in Haskell is used to abstract over the representation. The authors of [8] coined the term “symantic” to describe such a type class, because the interface of the type class defines the syntax of the DSL and the instances of the type class define the semantics. A separate type class instance is written for each representation, each corresponding to one member in the family of interpreters. This suits GOOL’s needs nicely, for GOOL requires multiple interpreters (one for each target language) of the same language (GOOL).

GOOL is a collection of these “symantic” type classes and instances. For organizational purposes, the functions comprising GOOL’s syntax are split across many type classes, roughly based on the internal types upon which the functions act. GOOL’s internal types correspond to the types of code being represented, not the types of the values in the target languages, though adding value-types to GOOL’s type system is planned for the future, discussed further in Section 7. Examples of internal types in GOOL, then, are Variable, Value, Type, Scope, Statement, Method, Class, Module, and Package. A Statement in GOOL, for example, is a representation of a piece of code that is a statement. Below is an excerpt of GOOL’s definition for the VariableSym type class; that is, the type class containing functions for defining variables in GOOL.

```
class (TypeSym repr)
=> VariableSym repr where
  type Variable repr
  var :: Label -> repr (Type repr)
      -> repr (Variable repr)
```

Since variables have types, a type for representing variables must also know how to represent types, thus we constrain our repr type variable to be an instance of TypeSym, the symantic type class containing functions for representing types, like the int, float, string, and listType functions described in Section ??.

The types used in the signatures of GOOL’s functions have the general form repr (X repr) for some type X, examples of which can be seen in the type signature for var. To understand these types, first remember that we will write different instances of this type class for each different language GOOL targets. That means the repr type variable will be instantiated with a type that represents code in a specific target language. In the repr (X repr) types, the type variable repr appears twice because there are two layers of abstraction: abstraction over the target language, handled by the outer repr, and abstraction over the underlying types to which GOOL’s types map, handled by the inner repr.

This abstraction over the underlying types that GOOL’s types represent happens because X repr is a type family. The

first line of the body of the VariableSym type class is type Variable repr. This line declares the type family Variable repr. Since the type is parameterized by repr, each instance we write of this type class can define the Variable family member differently. Usually different renderers will use the same underlying type for a given GOOL type, but the ability to change it on a per-target-language basis is useful for when a language requires storing more or less information than others.

The type signature for the var function says that var takes a label and a representation of a type and yields a representation of a variable. This should not be surprising given how we used var to define variables like var "foo" int in Section ??.

Label is simply a Haskell type synonym for String.

Shown below is the instance of the type class excerpt shown above for generating Java code.

```
instance VariableSym JavaCode where
  type Variable JavaCode = VarData
  var = varD
```

We have instantiated repr as JavaCode, which is a monad defined as:

```
newtype JavaCode a = JC {unJC :: a}
```

This may look like we are tagging our values with JC, in contradiction of the finally-tagless style, but Haskell’s compiler does not actually treat types defined with newtype as tags, so this is valid. unJC extracts the underlying value from a value wrapped in JavaCode. Calling unJC on a program written in GOOL allows Haskell to infer that the repr in the GOOL program should be concretized as JavaCode and to resolve the functions using the definitions from the JavaCode instances of the type classes, thereby generating Java code.

The second line of the instance states that the underlying type for Variables in Java is VarData, which is a record that holds data about the variable to be used later. The definition of VarData is shown below.

```
data VarData = VarD {
  varBind :: Binding ,
  varName :: String ,
  varType :: TypeData ,
  varDoc  :: Doc }
```

Stored in the VarData structure is a variable’s binding, either Static or Dynamic, the name of the variable as a String, the type as a TypeData structure, which is the underlying type for Types in GOOL for Java, and how the variable should appear in the generated code, represented as a Doc. Doc comes from the Text.PrettyPrint.HughesPJ Haskell package and represents formatted text. We use it to pretty-print our generated code, making it more readable. In common between all of GOOL’s underlying data structures is that they each contain a Doc. For some types, like a Block, the

underlying type is in fact nothing more than a Doc. The Package-level Docs are ultimately what will be printed to files to generate code. Keeping in mind that a VarData structure will usually be wrapped in a repr, such as JavaCode, the pieces of the wrapped VarData can be accessed in one of two ways. The first way is to use Haskell's fmap in combination with the built-in accessor functions for VarData. For example, fmap varDoc will take a VarData and return a JavaCode Doc. Functions on Docs can then be lifted to work on JavaCode Docs using Haskell's liftA family of functions. This method only works if the Variable repr is known to be equivalent to VarData, though. That is, it works in the context of JavaCode but not in a generic repr context. For when one needs to access a piece of VarData while in a generic repr context, GOOL offers its own “syntactic” accessor functions. For example, variableDoc is part of the VariableSym type class, with signature:

```
variableDoc :: repr (Variable repr)
              -> Doc
```

and definition for the JavaCode instance:

```
variableDoc = varDoc . unJC
```

Since variableDoc is part of the type class, it can be used in a generic repr context, and since its definition is in an instance of the type class where Variable repr is known to resolve to VarData, it can be defined simply by calling the varDoc accessor on the unwrapped VarData. GOOL offers many functions similar to variableDoc, for accessing the other pieces of VarData and for accessing the pieces of the underlying data structures for other internal GOOL types.

The underlying data types for different GOOL types each store specific information related to the type in question. Some examples are that the underlying type for Statements stores a Terminator which determines whether the statement should end in a semi-colon, the underlying type for Methods stores a Boolean for whether it is the main method, and the underlying types for Values, UnaryOps, and BinaryOps store precedence information so expressions can be printed without superfluous parentheses, leading to more readable code.

The third line of the VariableSym JavaCode instance defines the var function by dispatching to the varD function, defined as:

```
varD :: (RenderSym repr) => Label ->
      repr (Type repr) ->
      repr (Variable repr)
varD n t = varFromData Dynamic n t
          (varDocD n)

varDocD :: Label -> Doc
varDocD = text
```

varD constructs the variable using a function called varFromData, which is a GOOL function for constructing a variable in the generic repr context by explicitly passing the pieces needed to create a varData, as can be seen by its type signature:

```
varFromData :: Binding -> String ->
              repr (Type repr) -> Doc ->
              repr (Variable repr)
```

This function resides in a GOOL type class called InternalVariable. GOOL has a number of these “internal” type classes, none of which are exported as part of GOOL's user-facing interface. They contain functions that are useful for GOOL's various language renderers, but not to users. A user would not want to define a Variable by explicitly passing in the Doc, for example, which is why varFromData is not exposed. The Doc representation for the generated code is entirely internal to GOOL, not exposed in any of GOOL's user-facing functions. Examples of other GOOL functions that are only available internally are printSt for generating a low-level print statement, because it is superseded by higher-level functions for printing, and cast for casting between types, because any required type-casting is handled by higher-level functions.

In the definition for varD, the Doc for the variable is constructed simply by passing the variable name to text, which is a function from the HughesPJ package for converting a String to a Doc of that String. This should make sense because you can refer to a variable in Java simply by typing the variable's name. In fact, this is true of most OO languages, and this is why the functionality for var is defined in a generic varD function instead of defining it directly in the VariableSym JavaCode instance: the varD function can be re-used between language renderers. In the VariableSym instances for the other languages GOOL targets, the var function is defined as a dispatch to the varD function, exactly as it is for the Java instance. By writing generic functions for generating code that is common between target languages, we maximize code re-use and minimize effort required to write a renderer for a different language. The more similarities a language has to those that GOOL already targets, the less needs to be done to write the new renderer. GOOL's Java and C# renderers demonstrate this fact well. Out of 327 functions across all of GOOL's type classes, the instances of 227 of them are shared between the Java and C# renderers, in that they are just calls to the same common function. A further 37 are partially shared, for example they call the same common function but with different parameters. 143 functions are actually the same between all 4 languages GOOL currently targets, which might mean that some of them need not be included in the type class mechanism at all, and can instead be defined globally for all renderers, though this requires further investigation.

Examples from the Python and C# renderers are not shown here because they both work very similarly to the Java renderer. There are PythonCode and CSharpCode analogs to

JavaCode, the underlying types are all the same, and the functions are defined by calling common functions where possible or by constructing the GOOL value directly in the instance definition, if the definition is unique to that language.

C++ is different, however, because most modules are split between a source and header file. One module written in GOOL essentially needs to be traversed twice, once to generate the source file and a second time to generate the header file corresponding to the same module. To accomplish this, two instances of the GOOL type classes were written for C++ for two different types defined similarly to JavaCode: CppSrcCode for source code and CppHdrCode for header code. Since a main function does not require a header file, the CppHdrCode instance for a module containing only a main function is empty. The renderer is written such that an empty module or file does not actually get generated, and this is true of all GOOL's current renderers.

Since C++ source and header code should always be generated together, we needed to tie the CppSrcCode and CppHdrCode types together. To do this, a third type, CppCode, was created to pair the source and header types, defined as:

```
data CppCode x y a =
  CPPC {src :: x a, hdr :: y a}
```

The type variables x and y are intended to be instantiated with CppSrcCode and CppHdrCode, but they are left generic for now because we need to make an instance of a Pair type class for CppCode, as follows:

```
class Pair p where
  pfst :: p x y a -> x a
  psnd :: p x y b -> y b
  pair :: x a -> y a -> p x y a

instance Pair CppCode where
  pfst (CPPC xa _) = xa
  psnd (CPPC _ yb) = yb
  pair = CPPC
```

In summary, the Pair instance allows us to access the source code piece of the pair with pfst, the header code piece with psnd, or to combine a source and header piece to form a pair with pair. We then need to make CppCode instances of the GOOL type classes. The instance of VariableSym for the excerpt of the type class we have used as an example throughout this section is shown below.

```
instance (Pair p) => VariableSym
  (p CppSrcCode CppHdrCode) where
  type Variable (p
    CppSrcCode CppHdrCode) = VarData
  var n t = pair (var n $ pfst t)
    (var n $ psnd t)
```

Rather than write the instance specifically for CppCode, we wrote it for any Pair to make this code more re-usable. Unfortunately, the types making up the Pair had to be concretized as CppSrcCode and CppHdrCode so that the Haskell compiler could be sure that the VarData underlying type declared for the Pair matched the underlying types for each piece of the pair. The compiler knows that Variable CppSrcCode and Variable CppHdrCode are equivalent to VarData, but it cannot be sure that any generic Variable repr will be. The actual function definitions for these instances are trivial. The definition for var simply calls var in the context of CppSrcCode, calls it again in the context of CppHdrCode, and combines the results in a new pair. This same pattern is used for all of the function definitions for the Pair instances, so to write a new renderer for another language that requires two files per module, these Pair instances can be re-used just by changing the CppSrcCode and CppHdrCode types to the corresponding types for the new language renderer, and changing any underlying types that are different. This same technique can be used to write a renderer for a language that has even more than two files per module by writing an analog to the Pair type class that combines the appropriate number of constituents.

The last component of the C++ renderer is a function analogous to JavaCode's unJC, a function that can be called on a GOOL program to tell the Haskell compiler to concretize the repr as CppCode and therefore generate C++ code. The function is unCPPC, defined as:

```
unCPPC ::
  CppCode CppSrcCode CppHdrCode a -> a
unCPPC (CPPC (CPPSC a) _) = a
```

At the level of a Program, it no longer makes sense to consider source and header files independently, so the Pair instance of the prog function combines the source and header files together in the CppSrcCode piece of the Pair. This is why unCPPC simply extracts the CppSrcCode portion of the Pair—at that point, that portion contains both the source and header files. Note that the choice to combine them in the CppSrcCode half instead of the CppHdrCode half of the Pair was arbitrary.

After an unRepr function, like unJC or unCPPC, has been called on a GOOL program, the resulting pretty-printed Docs are transformed back into a Strings using the render function, and then written to files using standard Haskell IO.

5 Higher-level GOOL functions

GOOL provides many high-level functions that abstract over common OO patterns, allowing one to write a small amount of GOOL code to generate a large amount code in the target language, conforming to the target language's idioms, which may be very different between languages. This section shows examples of these high-level functions and the code they

generate, starting with those that abstract simple patterns on the scale of values, and working up to those that abstract more complex patterns on the scale of entire blocks of code or functions.

One of the simplest examples of a high-level function in GOOL is `sin`, the function for applying the trigonometric sine function to a value. Most OO languages do offer this function, though usually as part of a library and not built in to the language itself. If a GOOL user wanted to apply a sine function and GOOL did not offer a function for doing this, the user would have to use `extFuncApp` and pass it the library name and function name. But the library and function names are likely to be different between target languages, so to target multiple languages, they would have to change the GOOL code for each. A high-level `sin` function eliminates this problem. Each language renderer can define `sin` by calling the specific library and function for that language, but all the GOOL user needs to do is write `sin foo` for some declared variable `foo`. This will yield `math.sin(foo)` in Python, `Math.sin(foo)` in Java, `Math.Sin(foo)` in C#, and `sin(foo)` in C++. In addition to `sin`, GOOL offers functions for applying remaining trigonometric functions and other common functions in mathematics: `log`, `ln`, `exp`, `floor`, and `ceiling`.

Another simple pattern GOOL abstracts is accessing arguments passed to a main function through the command line. The name given to the variable holding the list of arguments is different depending on the target language, so GOOL offers the `argsList` function to represent the value of that variable. GOOL also has functions for certain actions commonly taken with that list, such as `arg` which, when passed an integer index, represents the value of the argument at that index, and `argExists` which represents a Boolean value for whether an argument exists at a given index.

Lists are a very common data structure in OO code, so GOOL provides a suite of functions for working with lists. For example, `listAccess` can be passed a list and index value and will return the value of the element of the list at the given index. The GOOL code `listAccess (valueOf ages) (litInt 1)` for the `ages` variable defined in Section ?? will generate `ages[1]` in Python and C#, `ages.get(1)` in Java, and `ages.at(1)` in C++. Other functions for working with lists offered by GOOL are `listSize` for getting the size of a list, `listAppend` for appending a value to a list, `listAdd` for adding a value to a list at a given index, `listSet` for changing the value of a list at a given index to a given value, `listIndexExists` for checking if a list has a value at a given index, and `indexOf` for getting the index of a given value in a given list. Slicing a list can be done with `listSlice` by passing it a variable to assign the sliced list to, a list to slice, and three values representing the starting and ending indices for the slice and the step size. The values are wrapped in Haskell's Maybe monad and default to the start of the list, the end of the list, and a step size of one if `Nothing` is passed. So

to take the elements from indices 1 and 2 of `ages` and assign them to a new list, `someAges`, the GOOL code looks like:

```
listSlice someAges (valueOf ages)
  (Just $ litInt 1) (Just $ litInt 3)
  Nothing
```

List slicing is of particular note because this action is much easier to do in Python than the other languages GOOL generates. The generated Python code for the above list slice is:

```
someAges = ages [1:3:]
```

While the generated Java code is shown below. Throughout this section, backslashes in generated code snippets indicate manually inserted line breaks, added so the code fits between the margins.

```
ArrayList<Double> temp = \
  new ArrayList<Double>(0);
for (int i_temp = 1; i_temp < 3; \
  i_temp++) {
  temp.add(ages.get(i_temp));
}
someAges = temp;
```

The generated C# and C++ code blocks are similar in structure to the Java code. This example demonstrates GOOL's idiomatic code generation. It would have been easy to generate Python code with the same structure as the other languages, and the code would even be correct. However, since list slicing is available as a high-level function in GOOL, the renderers have the freedom to define that function in a way best-suited to the target language, and an idiomatic, more readable version was generated in Python as a result.

GOOL also offers high-level functions for printing. `print` prints a value, `println` prints a value followed by a new line character, and `printFile` and `printFileLn` are the same for printing to a file. There are also variations that are shortcuts for when the value to print is just a literal string, such as `printFileStrLn`. Unlike print functions in many of the target languages, these print functions can be used on lists. So the GOOL code `println ages` generates `print(ages)` in Python, but in the other languages it generates code to loop through the list and print each element, like the Java code shown below.

```
System.out.print("[");
for (int list_i1 = 0; \
  list_i1 < ages.size() - 1; \
  list_i1++) {
  System.out.print(ages.get(list_i1));
  System.out.print(", ");
}
if (ages.size() > 0) {
  System.out.print(ages.get( \
```



```

881         ages.size() - 1));
882     }
883     System.out.println("");

```

This is another example where the renderers are clearly conforming to the idioms of a language. In addition to its functions for printing output, GOOL also offers functions for reading input, such as `getInput` and `getFileInput`.

Moving to larger-scale patterns, GOOL's `inOutFunc` generates a function based on three lists of variables: one of inputs, one of outputs, and one of variables that are both inputs and outputs. Consider a function `applyDiscount` that takes a price and a discount, subtracts the discount from the price, and returns both the new price and a Boolean for whether the price is below 20. It can be written in GOOL using `inOutFunc`, assuming variables `price`, `discount`, and `isAffordable` have been defined:

```

898 inOutFunc "applyDiscount" public static_
899     [discount] [isAffordable] [price]
900     (bodyStatements [
901         price &-= valueOf discount ,
902         isAffordable &=
903         valueOf price ?< litFloat 20.0])

```

The price is both an input and output, so it is in the third list. The discount is an input only and `isAffordable` is an output only, so they go in the first and second lists, respectively. This function has multiple outputs—price and `isAffordable`—and each of GOOL's target languages handles functions with multiple outputs differently. In Python, return statement with multiple values is used:

```

912 def applyDiscount(price , discount):
913     price = price - discount
914     isAffordable = price < 20
915
916     return price , isAffordable

```

In Java, the outputs are returned in an array of Objects:

```

920 public static Object[] applyDiscount( \
921     int price , int discount) \
922     throws Exception {
923     Boolean isAffordable;
924
925     price = price - discount;
926     isAffordable = price < 20;
927
928     Object[] outputs = new Object[2];
929     outputs[0] = price;
930     outputs[1] = isAffordable;
931     return outputs;
932 }

```

In C#, the outputs are passed as parameters, using the `out` keyword if it is only an output or the `ref` keyword if it is both an input and an output:

```

939 public static void applyDiscount( \
940     ref int price , int discount , \
941     out Boolean isAffordable) {
942     price = price - discount;
943     isAffordable = price < 20;
944 }

```

And in C++, the outputs are passed as pointer parameters:

```

947 void applyDiscount(int &price , \
948     int discount , bool &isAffordable) {
949     price = price - discount;
950     isAffordable = price < 20;
951 }

```

The structure of the function in each language is different, from the parameters to the function body to the return type. But each uses the language's most natural way of defining a function with multiple outputs. GOOL generates any needed variable declarations and return statements automatically, so the GOOL user is saved from typing these lines out manually. Functions defined with `inOutFunc` can be called with `inOutCall`, which again accepts the three lists of inputs, outputs, and those that are both. Through `inOutCall`, GOOL will again automatically generate any needed variable declarations and assignments, such as declaring the outputs array and then assigning its elements to the appropriate variables in Java.

In OO programming, it is common to write getter and setter methods in a class, so GOOL abstracts over these patterns as well. `getMethod` can be used to define a getter, and `setMethod` for a setter. Each needs only be provided the name of the class to which the method will belong, and the variable to be get or set. So, assuming the class `FooClass` has a variable `foo`, a getter for `foo` is simply `getMethod "FooClass" foo` and a setter is simply `setMethod "FooClass" foo`. The generated set method in Python looks like:

```

977 def setFoo(self , foo):
978     self.foo = foo

```

In Java:

```

980 public void setFoo(int foo) \
981     throws Exception {
982     this.foo = foo;
983 }

```

In C#:

```

986 public void setFoo(int foo) {
987     this.foo = foo;
988 }

```

And in C++:

```
void FooClass::setFoo(int foo) {
    this->foo = foo;
}
```

We show this example not only to demonstrate how a single, small line of GOOL code can generate much more code in the target language, but also to visualize some of the idiosyncracies present in the target languages. Even for such a simple function, there are subtle differences in each language that would be difficult to keep track of if someone were trying to write these programs manually, and GOOL saves the programmer from this tedium. For calling methods defined with `getMethod` or `setMethod`, GOOL also provides `get` and `set` functions, which must be passed the value of the object on which the method should be called, the variable to get or set, and, in the case of `set`, the value to set it to.

The final category of higher-order functions we will discuss are those that abstract over the design patterns described in [10]. GOOL currently has syntax for defining instances of three design patterns: Observer, State, and Strategy. The versions of these design patterns GOOL generates are simplified and small in scale. In the case of the Strategy pattern, Haskell does the work of storing and checking which strategy to use, only actually generating code for a strategy when it is used. `runStrategy` is the user-facing function for using the Strategy pattern in GOOL. It must be passed the name of the strategy to use, a list of pairs of strategy names and bodies, and Maybe a variable and value to assign after running the strategy. Haskell will check if the given strategy name is in the list, and simply generate the corresponding body if it is.

For the Observer pattern, `initObserverList` can be passed a list of values and will generate a declaration of an observer list variable initially containing the given values. `addObserver` can then be used to add a given value to the observer list, and `notifyObservers` will call a method on each of the observers. The name of the observer list variable is fixed, so there can only be one observer list in a given scope.

For the State pattern, `initState` will take a name and a state label and generate a declaration of a variable with the given name and assign it the given state. The states are just literal strings. `changeState` takes the variable name and a new state, and changes the state of that variable to the new state. And `checkState` takes the name of the state variable, a list of value-body pairs, and a fallback body, and generates a conditional (usually a switch statement) that checks the state and runs the corresponding body, or the fallback body if none of the states match.

The functionality granted by these high-level design pattern functions was already possible with GOOL's other functions. But they are useful because they are tailored to specific design patterns, so they are concise, for example by not requiring the user to manually define a loop for notifying

observers, and their syntax should be easily understood by those familiar with OO programming.

6 Related Work

We divide the Related Work into the following categories

- General-purpose code generation
- Multi-language OO code generation
- Design pattern modeling and code generation

which we present in turn.

Haxe is a general-purpose multi-paradigm language and cross-platform compiler. It compiles to all of the languages GOOL does, in addition to many others. However, it does not offer the high-level abstractions GOOL provides [3] (better reference?). Also, Haxe strips comments and generates source code around a custom framework which users would not be familiar with, so the generated code is not very readable.

Proton's 2nd version is a DSL and code generator for Java and C++, where the generator is designed to be capable of producing general-purpose imperative or object-oriented code. The Proton generator is model-driven and uses a final "output model" from which actual code can be trivially generated. Since the "output model" was so similar to the generated code, it presented challenges with regards to semantic, conventional, and library-related differences between the target language [13]. GOOL's finally-tagless approach and syntax for high-level tasks, on the other hand, helped it overcome differences between target languages.

ThingML [11] is a DSL for model-driven engineering targeting C, C++, Java, and JavaScript. While it can be used in a broad range of application domains, they all fall under the umbrella domain of distributed reactive systems, and so it is not quite a general-purpose DSL, unlike GOOL. ThingML's modelling-related syntax and abstractions are a contrast to GOOL's object-oriented syntax and abstractions. The generated code lacks some of the pretty-printing provided by GOOL, specifically indentation, which detracts from the readability.

Moving on to OO-specific code generators with multiple target languages, there are many examples of such DSLs, but for more restricted domains than GOOL. Google protocol buffers is a DSL for serializing structured data, which can then be compiled into Java, Python, Objective C, and C++ [2]. Thrift is a Facebook-developed tool for generating code in multiple languages and even multiple paradigms based on language-neutral descriptions of data types and interfaces [19]. Clearwater is an approach for implementing DSLs with multiple target languages for components of distributed systems [20]. The Time Weaver tool uses a multi-language code generator to generate "glue" code for real-time embedded systems [9]. The domain of mobile applications is host to a bevy of DSLs with multiple target languages, of which

MobDSL [14] and XIS-Mobile [17] are two examples. Conjure is a DSL for generating APIs. It reads YAML descriptions of APIs and can generate code in Java, TypeScript, Python, and Rust [1] (include this?). All of these are examples of multi-language code generation, but none of them generate general-purpose code like GOOL does.

A number of languages for modeling design patterns have been developed. The Design Pattern Modeling Language (DPML) [15] is similar to the Unified Modeling Language (UML) but designed specifically to overcome UML's shortcomings to be able to model all design patterns. DPML consists of both specification diagrams and instance diagrams for instantiations of design patterns, but does not attempt to generate actual source code from the models. The Role-Based Metamodeling Language [12] is also based on UML but with changes to allow for better models of design patterns, with specifications for the structure, interactions, and state-based behaviour in patterns. Again, source code generation is not attempted. Another metamodel for design patterns includes generation of Java code [4], and IBM developed a DSL for generation of OO code based on design patterns [6]. IBM's DSL was in the form of a visual user interface rather than a programming or modeling language. The languages that generate code do so only for design patterns, not for any general-purpose code like GOOL does.

7 Future Work

Currently GOOL code is typed based on what kind of code it represents: variable, value, type, or method, for example. Code that represents a variable or value is not further typed based on what the type of the variable or value would be in a traditional programming language: Boolean or integer, for example. There is nothing to stop a user from passing a non-list value to a function specifically intended for lists, like `listSize`, or from passing string values to numeric operators like `#+`. We plan on adding this additional layer of typing, making GOOL a statically typed programming language. We have started work to this end by making the underlying types for GOOL's Variables and Values Generalized Algebraic Data Types (GADTs), such as this one for Variables:

```
data TypedVar a where
  BVr :: VarData -> TypedVar Boolean
  IVr :: VarData -> TypedVar Integer
  ...
```

Members of the type family `Variable repr` would now map to `TypedVar` and the type for a `Variable` in the generic `repr` context would now be something like `repr (Variable repr Boolean)` instead of just `repr (Variable repr)`. All instances of `TypedVar` are built from the same `VarData` structure, but variables built with the different GADT constructors will have different types and Haskell's compiler will throw errors if a wrongly-typed variable is passed to a function.

There are many improvements we plan on making to the generated code, especially with regards to not generating code that is not necessary. For example, we currently always generate import statements for certain libraries, like `math` and `IO`-related libraries, but we'd instead like to detect when a library is used and only import those that are actually used. We plan on using Haskell's State monad to have a list of libraries that gets updated whenever a new library is used, and can then be read by the `buildModule` function to generate only the necessary imports. This technique of using the State monad will be useful for other improvements as well, such as to only generate `throws Exception` in the header of a Java method that actually may throw an exception, rather than in every method as is done currently.

We plan on adding syntax to interface with more kinds of external libraries, similar to how we currently interface with `math` libraries with functions like `sin`. We would like to be able to interface with libraries for solving Ordinary Differential Equations (ODEs), for instance, with different functions available for using different ODE-solving algorithms. Since interfacing with ODE-solving libraries can have major impact on the structure of a program, implementing these high-level functions may require having multiple passes over a GOOL program: an initial pass to collect information, such as the information that an ODE library is used, and then a second pass to generate the code.

Currently GOOL forces certain design decisions on the user, but we plan on providing configuration options to give the user more control over the code they generate. For example, GOOL uses `ArrayLists` to represent lists in Java, but there are many other list implementations that could be used, such as `Vector`. This is something we would like the user to have control over. Another example, building on the aforementioned plan to provide functions to interface with more external libraries, is to allow the user to choose which specific external library they want to be used. These kinds of choices could be accomplished by having the user pass a configuration (such as a Haskell record) to GOOL, which GOOL will then read to decide how the code should be generated.

GOOL's current set of high-level functions for generating common OO patterns is by no means exhaustive. We plan to continue to identify patterns for which GOOL can provide abstractions, and to implement those as more high-level functions, with the aim of writing OO programs in GOOL as efficiently as possible without losing expressivity.

8 Conclusion

OO programming languages are similar enough that it is feasible to have a single OO language that can be compiled to any other OO language. GOOL is a DSL with multiple target languages for the domain of OO GPLs. Like most DSLs, GOOL provides abstractions suited to its domain, in this case

abstractions of common OO patterns. Unlike most DSLs, GOOL considers the code it generates to be a product for human consumption in addition to computer consumption, so it focuses on generating idiomatic, human-readable, and documented code. The goal of generating idiomatic code is helped by the abstractions: GOOL provides syntax for expressing OO patterns naturally and efficiently, and these high-level functions afford each target language renderer the freedom to generate code following its own idioms. The goal of generating documented code is realized by GOOL's syntax for generating informal documentation in the form of code comments, or more formal documentation in the form of Doxygen-style structured comments for functions, classes, and modules. GOOL can even generate Doxygen configuration files and makefiles to facilitate compiling the documentation into PDF or HTML formats. The generated code is pretty-printed so that it is readable, and GOOL allows organization of code related to the same task into blocks to further increase readability. The idiomaticity and presence of documentation in the generated code also contribute to readability.

References

- [1] [n. d.]. Conjure: a code-generator for multi-language HTTP/JSON clients and servers. <https://palantir.github.io/conjure/#/> Accessed 2019-09-16.
- [2] [n. d.]. Google Protocol Buffers. <https://developers.google.com/protocol-buffers/> Accessed 2019-09-16.
- [3] [n. d.]. Haxe - The cross-platform toolkit. <https://haxe.org> Accessed 2019-09-13.
- [4] Hervé Albin-Amiot and Yann-Gaël Guéhéneuc. 2001. Meta-modeling design patterns: Application to pattern detection and code synthesis. In *Proceedings of ECOOP Workshop on Automating Object-Oriented Software Development Methods*.
- [5] Giora Alexandron, Michal Armoni, Michal Gordon, and David Harel. 2012. The effect of previous programming experience on the learning of scenario-based programming. In *Proceedings of the 12th Koli Calling International Conference on Computing Education Research*. ACM, 151–159.
- [6] Frank J. Budinsky, Marilyn A. Finnie, John M. Vlissides, and Patsy S. Yu. 1996. Automatic code generation from design patterns. *IBM systems Journal* 35, 2 (1996), 151–171.
- [7] Raymond PL Buse and Westley R Weimer. 2009. Learning a metric for code readability. *IEEE Transactions on Software Engineering* 36, 4 (2009), 546–558.
- [8] Jacques Carette, Oleg Kiselyov, and Chung-chieh Shan. 2009. Finally tagless, partially evaluated: Tagless staged interpreters for simpler typed languages. *Journal of Functional Programming* 19, 5 (2009), 509–543.
- [9] Dionisio de Niz and Raj Rajkumar. 2004. Glue code generation: Closing the loophole in model-based development. In *10th IEEE Real-Time and Embedded Technology and Applications Symposium (RTAS 2004). Workshop on Model-Driven Embedded Systems*. Citeseer.
- [10] Erich Gamma. 1995. *Design patterns: elements of reusable object-oriented software*. Pearson Education India.
- [11] Nicolas Harrand, Franck Fleurey, Brice Morin, and Knut Eilif Husa. 2016. Thingml: a language and code generation framework for heterogeneous targets. In *Proceedings of the ACM/IEEE 19th International Conference on Model Driven Engineering Languages and Systems*. ACM, 125–135.
- [12] Dae-Kyoo Kim, Robert France, Sudipto Ghosh, and Eunjee Song. 2003. A uml-based metamodeling language to specify design patterns. In *Proceedings of Workshop on Software Model Engineering (WiSME), at UML 2003*. Citeseer.
- [13] Gábor Kövesdán and László Lengyel. 2017. Multi-Platform Code Generation Supported by Domain-Specific Modeling. *International Journal of Information Technology and Computer Science* 9, 12 (2017), 11–18.
- [14] Dean Kramer, Tony Clark, and Samia Oussena. 2010. MobDSL: A Domain Specific Language for multiple mobile platform deployment. In *2010 IEEE International Conference on Networked Embedded Systems for Enterprise Applications*. IEEE, 1–7.
- [15] David Mapelsden, John Hosking, and John Grundy. 2002. Design pattern modelling and instantiation using DPML. In *Proceedings of the Fortieth International Conference on Tools Pacific: Objects for internet, mobile and embedded applications*. Australian Computer Society, Inc., 3–11.
- [16] Marjan Mernik, Jan Heering, and Anthony M Sloane. 2005. When and how to develop domain-specific languages. *ACM computing surveys (CSUR)* 37, 4 (2005), 316–344.
- [17] André Ribeiro and Alberto Rodrigues da Silva. 2014. Xis-mobile: A dsl for mobile applications. In *Proceedings of the 29th Annual ACM Symposium on Applied Computing*. ACM, 1316–1323.
- [18] Jean Scholtz and Susan Wiedenbeck. 1990. Learning second and subsequent programming languages: A problem of transfer. *International Journal of Human-Computer Interaction* 2, 1 (1990), 51–72.
- [19] Mark Slee, Aditya Agarwal, and Marc Kwiatkowski. 2007. Thrift: Scalable cross-language services implementation. *Facebook White Paper* 5, 8 (2007).
- [20] Galen S Swint, Calton Pu, Gueyoung Jung, Wenchang Yan, Younggyun Koh, Qinyi Wu, Charles Consel, Akhil Sahai, and Koichi Moriyama. 2005. Clearwater: extensible, flexible, modular code generation. In *Proceedings of the 20th IEEE/ACM international Conference on Automated software engineering*. ACM, 144–153.

A Appendix

Text of appendix ...